

Digital Assessment I

Name: AMAN KUMAR

Roll No: 21BDS0241

GitHub Link: https://github.com/Aman88097/EDA_21BD0241

Dataset Name: countymurders.csv

Dataset Link:

<https://raw.githubusercontent.com/salemprakash/EDA/main/Data/countymurders.csv>

Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
```

Step 1: Load the Dataset

```
# Load the dataset
url =
"https://raw.githubusercontent.com/salemprakash/EDA/main/Data/countymu
rders.csv"
df = pd.read_csv(url)

# Display structure
df.head()
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 37349 entries, 0 to 37348
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   rownames              37349 non-null  int64
1   arrests               36845 non-null  float64
```

2	countyid	37349	non-null	int64
3	density	37349	non-null	float64
4	popul	37349	non-null	int64
5	perc1019	37349	non-null	float64
6	perc2029	37349	non-null	float64
7	percblack	37349	non-null	float64
8	percmale	37349	non-null	float64
9	rpcincmaint	37346	non-null	float64
10	rpcpersinc	37346	non-null	float64
11	rpcunemins	37346	non-null	float64
12	year	37349	non-null	int64
13	murders	37349	non-null	int64
14	murdrate	37349	non-null	float64
15	arrestrate	36845	non-null	float64
16	statefips	37349	non-null	int64
17	countyfips	37349	non-null	int64
18	execs	37349	non-null	int64
19	lpopul	37349	non-null	float64
20	execrate	37349	non-null	float64

dtypes: float64(13), int64(8)
memory usage: 6.0 MB

Step 2: Data Cleaning

```
# Check missing values
print("Missing Values:\n", df.isnull().sum())

# Handle missing values
df.fillna(df.mean(numeric_only=True), inplace=True)

# Drop duplicates
df.drop_duplicates(inplace=True)

# Convert categorical columns if any
df = pd.get_dummies(df, drop_first=True)

# Confirm cleaning
df.info()
```

```
Missing Values:
  rownames      0
arrests      504
countyid      0
density       0
popul         0
perc1019      0
perc2029      0
percblack     0
```

```

percmale      0
rpcincmaint   3
rpcpersinc    3
rpcunemins    3
year          0
murders       0
murdrate      0
arrestrate    504
statefips     0
countyfips    0
execs         0
lpopul        0
execrate      0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 37349 entries, 0 to 37348
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   rownames              37349 non-null  int64
1   arrests               37349 non-null  float64
2   countyid              37349 non-null  int64
3   density               37349 non-null  float64
4   popul                 37349 non-null  int64
5   perc1019              37349 non-null  float64
6   perc2029              37349 non-null  float64
7   percblack             37349 non-null  float64
8   percmale              37349 non-null  float64
9   rpcincmaint           37349 non-null  float64
10  rpcpersinc            37349 non-null  float64
11  rpcunemins            37349 non-null  float64
12  year                  37349 non-null  int64
13  murders               37349 non-null  int64
14  murdrate              37349 non-null  float64
15  arrestrate            37349 non-null  float64
16  statefips             37349 non-null  int64
17  countyfips            37349 non-null  int64
18  execs                 37349 non-null  int64
19  lpopul                37349 non-null  float64
20  execrate              37349 non-null  float64
dtypes: float64(13), int64(8)
memory usage: 6.0 MB

```

Step 3: Data Handling

```

# Display basic statistics
print(df.describe())

```

```

# Check correlation
print(df.corr())

# Detect outliers using Interquartile Range (IQR)
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1

# Remove outliers
df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 *
IQR))).any(axis=1)]
print("Data after removing outliers:", df.shape)

```

	rownames	arrests	countyid	density
popul \				
count	37349.000000	37349.000000	37349.000000	37349.000000
	3.734900e+04			
mean	18675.000000	6.782250	32921.927173	252.241067
	8.934354e+04			
std	10781.871939	49.789272	15528.352966	1663.768484
	2.718545e+05			
min	1.000000	0.000000	1001.000000	0.050000
	8.500000e+01			
25%	9338.000000	0.000000	20105.000000	17.681580
	1.314400e+04			
50%	18675.000000	1.000000	36065.000000	44.240000
	2.879200e+04			
75%	28012.000000	3.000000	48049.000000	106.600000
	6.648000e+04			
max	37349.000000	2391.000000	56045.000000	54058.770000
	9.127751e+06			

	perc1019	perc2029	percblack	percmale
rpcincmaint \				
count	37349.000000	37349.000000	37349.000000	37349.000000
	37349.000000			
mean	15.582640	14.584615	7.823194	43.350958
	165.450844			
std	1.973399	3.696407	13.287067	3.717612
	97.485046			
min	7.080000	5.617244	0.000000	35.150000
	5.490000			
25%	14.320000	12.301700	0.200000	40.900000
	96.250000			
50%	15.420000	14.270000	1.450000	41.810000
	145.170000			
75%	16.730000	16.200000	8.740000	45.870000
	209.880000			
max	30.484580	40.520000	86.279340	78.040000
	1306.496000			

	...	rpcunemins	year	murders	murdrate	\
count	...	37349.000000	37349.000000	37349.000000	37349.000000	
mean	...	70.557953	1988.000000	7.286915	0.508202	
std	...	52.907268	4.899045	47.217586	0.851044	
min	...	0.000000	1980.000000	0.000000	0.000000	
25%	...	35.200000	1984.000000	0.000000	0.000000	
50%	...	57.100000	1988.000000	1.000000	0.241044	
75%	...	89.958000	1992.000000	3.000000	0.735294	
max	...	642.730000	1996.000000	1944.000000	39.840640	

	arrestrate	statefips	countyfips	execs
lpopul	\			
count	37349.000000	37349.000000	37349.000000	37349.000000
mean	0.511486	32.821575	100.352299	0.006854
std	1.224217	15.503684	107.942699	0.112448
min	0.000000	1.000000	1.000000	0.000000
25%	0.000000	20.000000	33.000000	0.000000
50%	0.177154	36.000000	75.000000	0.000000
75%	0.691197	48.000000	127.000000	0.000000
max	148.658400	56.000000	840.000000	7.000000

	execrate
count	37349.000000
mean	0.001042
std	0.029068
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	2.388916

[8 rows x 21 columns]

	rownames	arrests	countyid	density	popul
perc1019	\				
rownames	1.000000	-0.042387	9.874345e-01	0.002649	-0.060608
arrests	-0.042387	1.000000	-4.581712e-02	0.311418	0.791257
countyid	0.987434	-0.045817	1.000000e+00	0.010010	-0.061212
density	0.002649	0.311418	1.000971e-02	1.000000	0.380315

0.135204						
popul	-0.060608	0.791257	-6.121241e-02	0.380315	1.000000	-
0.119220						
perc1019	0.039989	-0.054905	3.612727e-02	-0.135204	-0.119220	
1.000000						
perc2029	-0.005638	0.097409	-7.710277e-03	0.126121	0.185030	
0.256169						
percblack	-0.013782	0.118160	-2.857313e-02	0.129959	0.084739	
0.122132						
percmale	0.001813	-0.020178	-1.443695e-03	-0.020912	-0.013728	-
0.208278						
rpcincmaint	-0.037361	0.142790	-5.677194e-02	0.164865	0.099991	
0.098407						
rpcpersinc	-0.051810	0.133192	-4.822744e-02	0.270458	0.331269	-
0.462651						
rpcunemins	-0.078840	0.025265	-7.233701e-02	0.034143	0.044673	
0.066413						
year	0.000454	0.001576	3.461751e-13	0.003468	0.014997	-
0.382645						
murders	-0.039936	0.905295	-4.245941e-02	0.433851	0.877110	-
0.073315						
murdrate	0.030421	0.170679	1.384874e-02	0.137594	0.118712	
0.046229						
arrestrate	0.021427	0.154936	5.991958e-03	0.061239	0.064456	
0.036582						
statefips	0.987121	-0.045881	9.999771e-01	0.009837	-0.060965	
0.035966						
countyfips	0.270752	-0.001236	2.318303e-01	0.027039	-0.049541	
0.031413						
execs	0.020697	0.083565	1.865559e-02	0.022459	0.160758	-
0.019806						
lpopul	-0.046460	0.295137	-3.736177e-02	0.280124	0.557064	-
0.078025						
execrate	0.011430	-0.000748	9.194172e-03	-0.000814	-0.003113	-
0.011402						

	perc2029	percblack	percmale	rpcincmaint	...
rpcunemins \					
rownames	-0.005638	-0.013782	0.001813	-0.037361	... -
0.078840					
arrests	0.097409	0.118160	-0.020178	0.142790	...
0.025265					
countyid	-0.007710	-0.028573	-0.001444	-0.056772	... -
0.072337					
density	0.126121	0.129959	-0.020912	0.164865	...
0.034143					
popul	0.185030	0.084739	-0.013728	0.099991	...
0.044673					
perc1019	0.256169	0.122132	-0.208278	0.098407	...

0.066413						
perc2029	1.000000	0.174507	-0.223615	-0.146472	...	
0.032846						
percblack	0.174507	1.000000	-0.104799	0.391585	...	-
0.129341						
percmale	-0.223615	-0.104799	1.000000	0.113600	...	-
0.088447						
rpcincmaint	-0.146472	0.391585	0.113600	1.000000	...	
0.189517						
rpcpersinc	-0.006004	-0.130672	0.223169	-0.279905	...	-
0.082969						
rpcunemins	0.032846	-0.129341	-0.088447	0.189517	...	
1.000000						
year	-0.428285	0.011903	0.660332	0.245385	...	-
0.202772						
murders	0.117278	0.139215	-0.025080	0.157738	...	
0.025989						
murdrate	0.074479	0.312546	-0.064556	0.199867	...	-
0.024076						
arrestrate	0.042012	0.225487	-0.037317	0.151704	...	-
0.030825						
statefips	-0.008012	-0.029168	-0.001293	-0.056677	...	-
0.071220						
countyfips	0.041639	0.078872	-0.021957	-0.026593	...	-
0.177027						
execs	0.026597	0.047100	0.028234	0.005941	...	-
0.018663						
lpopul	0.424437	0.148421	-0.082901	0.056214	...	
0.127596						
execrate	-0.011871	0.020806	0.029087	0.007829	...	-
0.014464						

		year	murders	murdrate	arrestrate	
statefips \						
rownames	4.543780e-04	-0.039936	0.030421	0.021427	9.871205e-01	
arrests	1.575601e-03	0.905295	0.170679	0.154936	-4.588142e-02	
countyid	3.461751e-13	-0.042459	0.013849	0.005992	9.999771e-01	
density	3.468035e-03	0.433851	0.137594	0.061239	9.837386e-03	
popul	1.499733e-02	0.877110	0.118712	0.064456	-6.096488e-02	
perc1019	-3.826454e-01	-0.073315	0.046229	0.036582	3.596605e-02	
perc2029	-4.282845e-01	0.117278	0.074479	0.042012	-8.012452e-03	
percblack	1.190266e-02	0.139215	0.312546	0.225487	-2.916773e-02	

02					
percmales	6.603317e-01	-0.025080	-0.064556	-0.037317	-1.293116e-03
rpcincmaint	2.453846e-01	0.157738	0.199867	0.151704	-5.667713e-02
rpcpersinc	2.927135e-01	0.170730	-0.035015	-0.040336	-4.823532e-02
rpcunemins	-2.027722e-01	0.025989	-0.024076	-0.030825	-7.121958e-02
year	1.000000e+00	0.003815	-0.048111	-0.024652	3.456787e-13
murders	3.814626e-03	1.000000	0.212655	0.119755	-4.247505e-02
murdrate	-4.811082e-02	0.212655	1.000000	0.448823	1.318676e-02
arrestrate	-2.465185e-02	0.119755	0.448823	1.000000	5.431574e-03
statefips	3.456787e-13	-0.042475	0.013187	0.005432	1.000000e+00
countyfips	1.496656e-13	-0.007458	0.098245	0.081857	2.252368e-01
execs	4.452117e-02	0.155903	0.042837	0.017847	1.844054e-02
lpopul	1.689796e-02	0.343957	0.068035	0.034184	-3.678215e-02
execrate	3.038773e-02	-0.000010	0.005167	0.004100	9.045622e-03

	countyfips	execs	lpopul	execrate
rownames	2.707525e-01	0.020697	-0.046460	0.011430
arrests	-1.235829e-03	0.083565	0.295137	-0.000748
countyid	2.318303e-01	0.018656	-0.037362	0.009194
density	2.703860e-02	0.022459	0.280124	-0.000814
popul	-4.954060e-02	0.160758	0.557064	-0.003113
perc1019	3.141287e-02	-0.019806	-0.078025	-0.011402
perc2029	4.163893e-02	0.026597	0.424437	-0.011871
percblack	7.887198e-02	0.047100	0.148421	0.020806
percmales	-2.195738e-02	0.028234	-0.082901	0.029087
rpcincmaint	-2.659270e-02	0.005941	0.056214	0.007829
rpcpersinc	-9.889311e-03	0.050651	0.376050	-0.000379
rpcunemins	-1.770272e-01	-0.018663	0.127596	-0.014464
year	1.496656e-13	0.044521	0.016898	0.030388
murders	-7.457649e-03	0.155903	0.343957	-0.000010
murdrate	9.824459e-02	0.042837	0.068035	0.005167
arrestrate	8.185673e-02	0.017847	0.034184	0.004100
statefips	2.252368e-01	0.018441	-0.036782	0.009046
countyfips	1.000000e+00	0.035151	-0.091789	0.023437
execs	3.515088e-02	1.000000	0.096420	0.350218
lpopul	-9.178859e-02	0.096420	1.000000	-0.009481


```
execrate      2.343731e-02  0.350218 -0.009481  1.000000
```

```
[21 rows x 21 columns]
```

```
Data after removing outliers: (20206, 21)
```

Step 4: 1D, 2D, and N-D Visualization

```
# 1D: Histogram
```

```
df.hist(figsize=(12, 8))
```

```
plt.suptitle("1D Histograms")
```

```
plt.show()
```

```
# 2D: Correlation Heatmap
```

```
plt.figure(figsize=(10, 6))
```

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

```
plt.title("2D Heatmap")
```

```
plt.show()
```

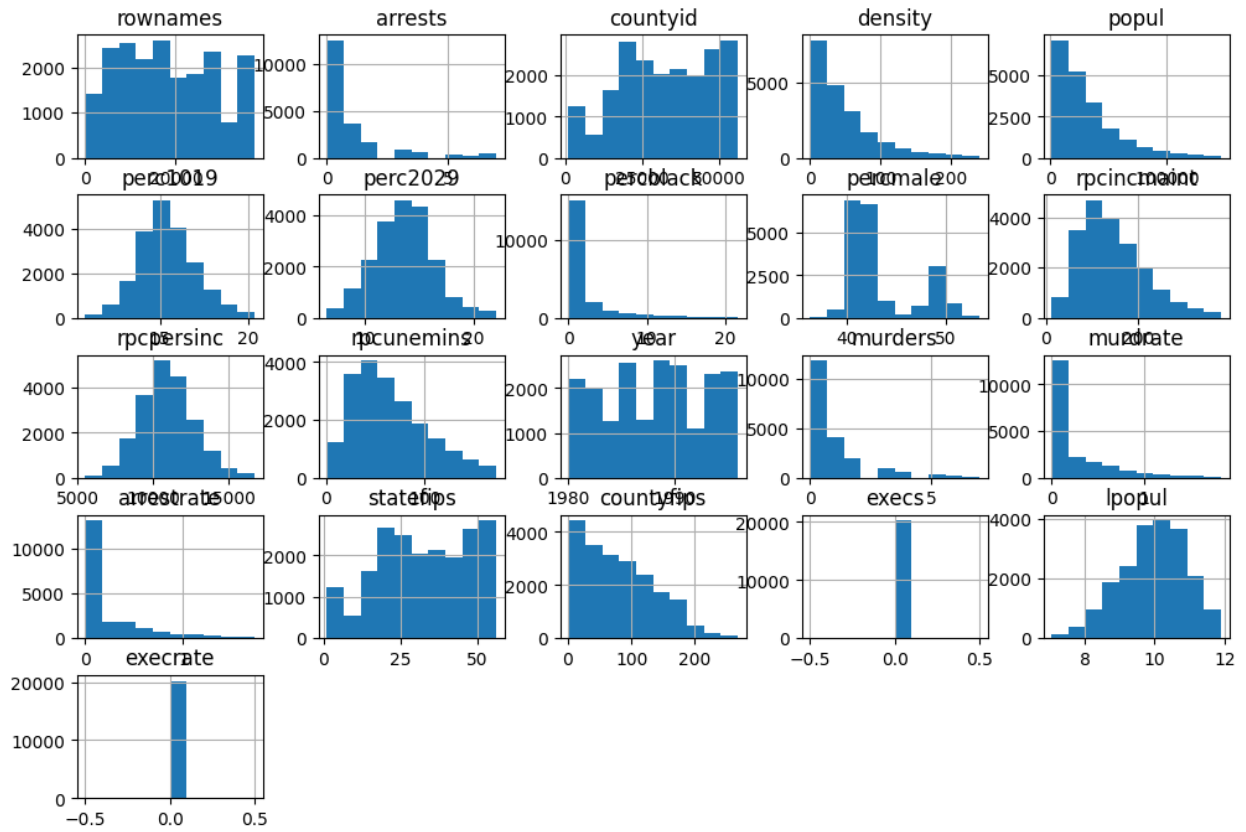
```
# N-D: Pairplot
```

```
sns.pairplot(df)
```


```
plt.suptitle("N-D Pairplot", y=1.02)
```


```
plt.show()
```

1D Histograms



	rownames	arrests	countyid	density	popul	perc1019	perc2029	percblack	percmale	rpcincmaint	rpcpersinc	rpcunemins	year	murders	murdrate	arrestrate	statefips	countyfips	execs	lpopul	execrate
rownames	1	0.014	0.99	0.038	0.038	0.028	0.053	0.007	0.120	0.053	0.138	0.009	0.250	0.081	0.078	0.99	0.014	0.038			
arrests	0.014	1	0.002	0.3	0.39	0.073	0.220	0.280	0.065	0.100	0.19	0.083	0.620	0.42	0.0	0.002	0.2023				
countyid	0.99	0.002	1	0.046	0.038	0.042	0.029	0.043	0.092	0.110	0.040	0.160	0.032	0.40	0.068	0.06	1	0.0018			
density	0.0340	0.30	0.046	1	0.7	0.012	0.350	0.260	0.027	0.035	0.240	0.081	0.034	0.39	0.070	0.080	0.46	0.091			
popul	0.033	0.39	0.038	0.76	1	0.065	0.38	0.20	0.021	0.060	0.240	0.160	0.043	0.530	0.11	0.10	0.038	0.078			
perc1019	0.038	0.078	0.042	0.012	0.065	1	0.39	0.018	0.13	0.013	0.33	0.11	0.30	0.049	0.046	0.42	0.032				
perc2029	0.028	0.220	0.029	0.350	0.380	0.39	1	0.16	0.41	0.170	0.170	0.12	0.55	0.230	0.11	0.10	0.29	0.034			
percblack	0.055	0.280	0.048	0.26	0.20	0.018	0.16	1	0.015	0.092	0.038	0.13	0.048	0.330	0.240	0.048	0.043				
percmale	0.007	0.065	0.009	0.020	0.021	0.180	0.41	0.018	1	0.210	0.240	0.083	0.71	0.078	0.076	0.48	0.091	0.014			
rpcincmaint	-0.12	0.051	0.110	0.035	0.060	0.019	0.170	0.092	0.21	1	-0.35	0.310	0.270	0.092	0.110	0.120	0.110	0.067			
rpcpersinc	-0.053	0.008	0.104	0.240	0.240	0.330	0.170	0.038	0.240	0.35	1	-0.21	0.380	0.042	0.093	0.098	0.470	0.032			
rpcunemins	-0.018	0.019	0.016	0.081	0.160	0.110	0.120	0.130	0.083	0.31	0.21	1	-0.190	0.050	0.140	0.061	0.13				
year	-0.005	0.283	0.003	0.30	0.043	0.380	0.550	0.004	0.71	0.270	0.380	0.19	1	-0.040	0.072	0.060	0.032	0.01			
murders	0.057	0.62	0.047	0.390	0.530	0.049	0.230	0.330	0.073	0.090	0.042	0.050	0.41	0.72	0.490	0.047	0.031				
murdrate	0.081	0.42	0.068	0.070	0.110	0.047	0.110	0.240	0.076	0.110	0.093	0.140	0.072	1	0.630	0.068	0.013				
arrestrate	0.078	0.7	0.060	0.0810	0.10	0.0460	0.1	0.260	0.048	0.120	0.093	0.046	0.060	0.490	0.63	1	0.060	0.018			
statefips	0.99	0.002	1	0.046	0.038	0.042	0.029	0.043	0.092	0.110	0.040	0.160	0.032	0.40	0.068	0.06	1	0.0021			
countyfips	0.018	0.023	0.001	0.091	0.078	0.032	0.034	0.043	0.018	0.067	0.032	0.130	0.010	0.034	0.018	0.018	0.002	1			
execs																					





Step 5: Time-Series Analysis

Group data by 'year' and calculate the sum or mean of relevant metrics

```
yearly_data = df.groupby('year').agg({
    'murders': 'sum',
    'arrests': 'sum',
    'arrestrate': 'mean',
    'murdrate': 'mean',
    'execrate': 'mean'
})
```


```

}))

# Plot time series trends
plt.figure(figsize=(12, 6))
for col in yearly_data.columns:
    plt.plot(yearly_data.index, yearly_data[col], marker='o',
             label=col)

plt.title("Time Series Analysis of Crime Trends Over Years")
plt.xlabel("Year")
plt.ylabel("Count / Rate")
plt.legend()
plt.grid(True)
plt.show()

```



Step 6: Dimensionality Reduction (PCA)

```

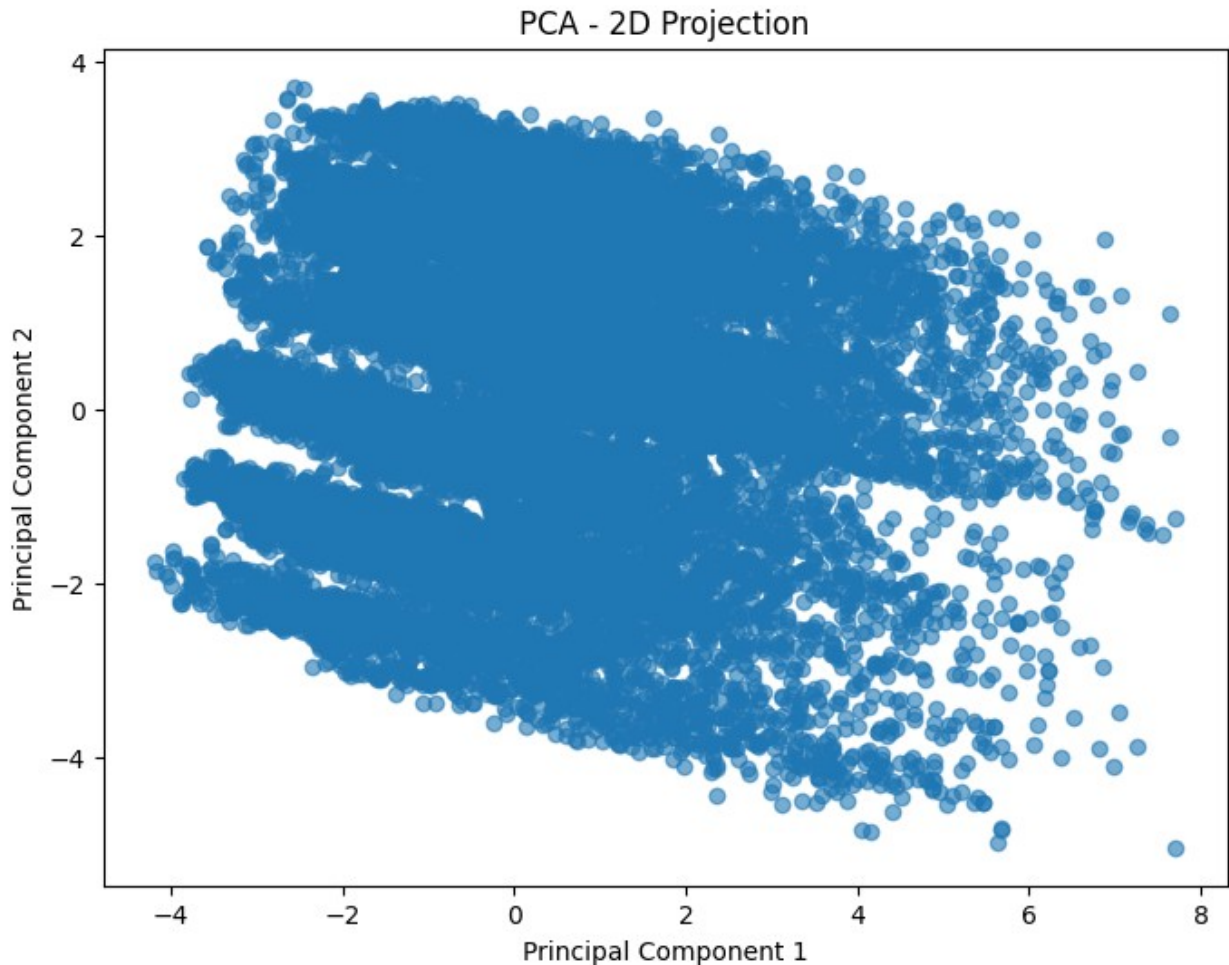
# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df)

# Apply PCA
pca = PCA(n_components=2)
pca_result = pca.fit_transform(X_scaled)

# PCA Visualization
plt.figure(figsize=(8, 6))

```

```
plt.scatter(pca_result[:, 0], pca_result[:, 1], alpha=0.6)
plt.title("PCA - 2D Projection")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.show()
```



Step 7: Model Building

```
# Create a binary target: high murder rate (above median)
df['high_murder_rate'] = (df['murdrate'] >
df['murdrate'].median()).astype(int)

# Drop columns not useful for modeling
drop_cols = ['rownames', 'murdrate', 'statefips', 'countyfips'] #
also dropping the original murdrate
X = df.drop(columns=drop_cols + ['high_murder_rate'])
y = df['high_murder_rate']
```

```

# Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42)

# Train a Random Forest with reduced complexity
model = RandomForestClassifier(n_estimators=50, max_depth=10,
random_state=42)
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

```

Step 8: Model Evaluation

```

# Evaluation
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test,
y_pred))

```

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2354
1	1.00	1.00	1.00	1688
accuracy			1.00	4042
macro avg	1.00	1.00	1.00	4042
weighted avg	1.00	1.00	1.00	4042